

Search for functions, topics, examples, tutorials... 🔍

CUDALink

CUDALINK TUTORIAL   Tutorials   Related Guides

Introduction

Getting Started

CUDA and CUDALink

CUDALink Application Areas

CUDALink Programming

CUDALink allows *Mathematica* to use the CUDA parallel computing architecture on Graphical Processing Units (GPUs). It contains functions that use CUDA-enabled GPUs to boost performance in a number of areas, such as linear algebra, financial simulation, and image processing. *CUDALink* also integrates CUDA with existing *Mathematica* development tools, allowing a high degree of automation and control.

Getting Started

To use any *CUDALink* functions, the application has to be loaded.

```
In[1]:= Needs["CUDALink`"]
```

CUDAQ tells you whether a CUDA-capable device is available and can be used.

```
In[2]:= CUDAQ[]
```

```
Out[2]= True
```

If CUDAQ returns `False`, then *CUDALink* will not work. For more information, read "[CUDALink Setup](#)".

CUDAInformation tells you more information about the graphics processing unit. Here the GPU is a Quadro FX unit with 96 cores and 1GB of graphics memory.

```
In[3]:= CUDAInformation[]
```

```
Out[3]=
```

```
{1 -> {Name -> Quadro FX 2800M, Clock Rate -> 1500000, Compute Capabilities -> 1.1,
GPU Overlap -> 1, Maximum Block Dimensions -> {512, 512, 64},
Maximum Grid Dimensions -> {65535, 65535, 1},
Maximum Threads Per Block -> 512, Maximum Shared Memory Per Block -> 16384,
Total Constant Memory -> 65536, Warp Size -> 32, Maximum Pitch -> 2147483647,
Maximum Registers Per Block -> 8192, Texture Alignment -> 256,
Multiprocessor Count -> 12, Core Count -> 96, Execution Timeout -> 1,
Integrated -> False, Can Map Host Memory -> False, Compute Mode -> Default,
Texture1D Width -> 8192, Texture2D Width -> 65536, Texture2D Height -> 32768,
Texture3D Width -> 2048, Texture3D Height -> 2048, Texture3D Depth -> 2048,
Texture2D Array Width -> 8192, Texture2D Array Height -> 8192,
Texture2D Array Slices -> 512, Surface Alignment -> 256,
Concurrent Kernels -> False, ECC Enabled -> False, Total Memory -> 1048117248}}
```

In *Mathematica* you can make matrices of real random numbers with the `RandomReal` command.

```
In[4]:= randM = RandomReal[1, {4000, 4000}];
```

You can multiply the matrix with itself, as shown below.

```
In[5]:= AbsoluteTiming[randM.randM];
```

```
Out[5]= {3.8204592, Null}
```

You can use the graphics processor to do the multiplication using `CUDADot`.

```
In[6]:= AbsoluteTiming[CUDADot[randM, randM];]
```

reference.wolfram.com/mathematica/CUDALink/tutorial/Introduction.html#1111406551

1/5

```
In[6]:= AbsoluteTiming[CUDADot[randM, randM];]
```

```
Out[6]:= {1.0840260, Null}
```

It is even faster to load the data onto the GPU with `CUDAMemoryLoad`. The result is a `CUDAMemory` expression; this can be used as a handle to the data for more computations.

```
In[7]:= randMG = CUDAMemoryLoad[randM]
```

```
Out[7]:= CUDAMemory[<10116>, Double]
```

Now you can pass the `CUDAMemory` to `CUDADot`, which stores the result on the GPU and returns a new `CUDAMemory` expression.

```
In[8]:= AbsoluteTiming[res = CUDADot[randMG, randMG]]
```

```
Out[8]:= {0.1357474, CUDAMemory[<16579>, Double]}
```


You can retrieve the data from the graphics processor with `CUDAMemoryGet`. Here, the dimensions of the result are shown to be as expected.

```
In[9]:= Dimensions[CUDAMemoryGet[res]]
```

```
Out[9]:= {4000, 4000}
```

## CUDA and CUDALink

CUDA is used as the computing engine for NVIDIA graphics processing units (GPUs); it provides a programming interface that can be called by software applications such as *Mathematica* with *CUDALink*. It allows CUDA GPUs to be used for parallel computations, allowing many concurrent threads to run. The following graph shows the performance of CPUs and GPUs—measured in billions of floating-point operations per second (GFLOP/s).



Evolution of CPU and GPU GFLOPS since 2003 (source: *NVIDIA CUDA Manual*).

*CUDALink* allows *Mathematica* users to call the CUDA programming layer directly; it also provides users higher-level functions, provided in a number of CUDA libraries, for solutions in areas such as high-performance core linear algebra and Fourier transforms.

## CUDALink Application Areas

*CUDALink* provides functions in various application areas. These include carefully tuned linear algebra, discrete Fourier transforms, and image processing algorithms. This section gives an introduction to some of these applications.

### Image Processing

*CUDALink* offers many image processing algorithms that have been carefully tuned to run on GPUs. These include the binary image operations (`CUDAImageAdd`, `CUDAImageSubtract`, `CUDAImageMultiply`, and `CUDAImageDivide`), the morphology operators (`CUDAERosion`, `CUDADilation`, `CUDAOpening`, and `CUDAClosing`), and image convolution (`CUDAImageConvolve`).

To use any of the *CUDALink* functionality, you first need to include the *CUDALink* application, as shown below.

```
In[1]:= Needs["CUDALink`"]
```

Now you can apply CUDA-based image processing functions directly to images. This example carries out channel-wise multiplication of two input images using `CUDAImageMultiply`.

```
In[2]:=
```

```
CUDAImageMultiply[ ,  ]
```

Out[2]=




Since the CUDA functions are *Mathematica* functions, they can be used in conjunction with other functions (like `Manipulate`). In the following example, a linear interpolation is carried out with `Manipulate` being used to vary the interpolation parameter.

In[3]:=

```
Manipulate[CUDAImageAdd[CUDAImageMultiply[ , 1 - alpha],
```

```
CUDAImageMultiply[ , alpha]], {alpha, 0.0, 1.0, .001}]
```

Out[3]=



CUDA functions can also be used with *Mathematica*'s curated data, `Import` and `Export` functions, as well as its visualization functions. In the following, the core computation is done with CUDA, while using *Mathematica* for all the other functions.

```
In[4]:= gisData = Import["http://exampledata.wolfram.com/sdtsdem.tar.gz", "Data"];
ReliefPlot[CUDAImageConvolve[gisData, {{1, 0, -1}}],
ColorFunction -> "GreenBrownTerrain"]
```

Out[5]=



## Fourier Analysis

The functions `CUDAFourier` and `CUDAInverseFourier` carry out Fourier transforms and inverse

Fourier transforms using CUDA.

To use any of the *CUDALink* functionality, you first need to include the *CUDALink* application, as shown below.

```
In[1]:= Needs["CUDALink`"]

In[2]:= vec = Range[1., 10];
         CUDAFourier[vec]

Out[3]:= {17.3925 + 0. i, -1.58114 - 4.86624 i, -1.58114 - 2.17625 i, -1.58114 - 1.14876 i,
          -1.58114 - 0.513743 i, -1.58114 + 8.09721 × 10-16 i, -1.58114 + 0.513743 i,
          -1.58114 + 1.14876 i, -1.58114 + 2.17625 i, -1.58114 + 4.86624 i}
```

If the input to a CUDA function is a *CUDAMemory* handle, the result will also be a *CUDAMemory* handle.

```
In[4]:= gpuVec = CUDAMemoryLoad[vec];
         CUDAFourier[gpuVec]

Out[5]:= CUDAMemory[<7091>, ComplexDouble]
```

You can then retrieve the data from the GPU with *CUDAMemoryGet*.

```
In[6]:= CUDAMemoryGet[%]

Out[6]:= {17.3925 + 0. i, -1.58114 - 4.86624 i, -1.58114 - 2.17625 i, -1.58114 - 1.14876 i,
          -1.58114 - 0.513743 i, -1.58114 + 8.09721 × 10-16 i, -1.58114 + 0.513743 i,
          -1.58114 + 1.14876 i, -1.58114 + 2.17625 i, -1.58114 + 4.86624 i}
```

This general principle applies to all *CUDALink* functions. It makes it easy to test and develop, and then to work by keeping data on the GPU, which improves efficiency.

## CUDALink Programming

A key feature of *CUDALink* is how it makes it easy to develop new GPU programs and integrate them into your *Mathematica* work. This requires that you have installed a C compiler.

To use any of the *CUDALink* functionality, you first need to include the *CUDALink* application, as shown below.

```
In[14]:= Needs["CUDALink`"]
```

Here, a simple CUDA function is loaded. It takes an input vector and doubles each element.

```
In[15]:= doubleFun = CUDAFunctionLoad["
    __global__ void doubleVec(mint * in,  mint length) {
        mint index = threadIdx.x + blockIdx.x * blockDim.x;

        if (index < length)
            in[index] = 2*in[index];
    }", "doubleVec", {{_Integer}, _Integer}, 256]

Out[15]:= CUDAFunction[<>, doubleVec, {{_Integer}, _Integer}]
```

*CUDAFunctionLoad* requires that you have a C compiler installed. If this does not work for you, try to consult "C Compiler".

Here is an input vector that will be used to call the function.

```
In[16]:= vec = Range[20];
```

This calls the *CUDAFunction*.

```
In[17]:= doubleFun[vec, 20]

Out[17]:= {{2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30, 32, 34, 36, 38, 40}}
```

This loads data to the GPU and calls your function on the data. The result is a *CUDAMemory* expression.

```
In[5]:= gpuVec = CUDAMemoryLoad[vec];
         doubleFun[gpuVec, 20]

Out[6]:= {CUDAMemory[<17648>, Integer]}
```

This retrieves the result from the GPU.

```
In[7]:= CUDAMemoryGet[First[%]]

Out[7]:= {2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30, 32, 34, 36, 38, 40}
```

You can get information on your function with *CUDAFunctionInformation*, as shown in the following.

```
In[18]:= CUDAFunctionInformation[doubleFun]
```

```

Out[18]= {Source →
__global__ void doubleVec(mint * in,  mint length) {
    mint index = threadIdx.x + blockIdx.x * blockDim.x;

    if (index < length)
        in[index] = 2*in[index];
}, ID → 21362576, BinaryFile →
/home/usr2/dcampbell/.Mathematica/ApplicationData/CUDALink/BuildFolder/
dcampbell1x-29298/CUDAFUNCTION-959.cubin,
Function → doubleVec, ArgumentTypes → {_Integer}, _Integer},
Options → {Platform → 1, Device → 1, Defines → ,
SystemDefines → -DUSING_CUDA_FUNCTION=1 -Dmint=int -DReal_t=float,
SystemIncludeDirectories → , IncludeDirectories → , CompileOptions → ,
ShellcommandFunction → None, ShellOutputFunction → None,
TargetPrecision → Single, BlockDimensions → 256}}

```

### Related Guides

[Running Computations on GPU Using CUDA](#)  
[Running Computations on GPU Using OpenCL](#)  
[C/C++ Language Interface](#)  
[Calling C Compilers from \*Mathematica\*](#)

### Related Tutorials

[CUDALink User Guide](#)  
[Setup and Operations](#)  
[Functions](#)  
[Programming](#)  
[Memory](#)  
[Applications](#)  
[Multiple Devices](#)  
[Reference](#)

[Give Feedback](#)

#### **Mathematica 9 is now available!**

400+ new features, including the new Wolfram Predictive Interface, social network analysis, enterprise CDF deployment, and more »

[Try Now](#) [Buy/Upgrade](#)

#### **New to Mathematica?**

[Find your learning path »](#)

#### **Have a question?**

[Ask support »](#)